

Momentum Factor Construction and Signal Orthogonality: A Mathematical Framework for Systematic Equity Strategies

Gilberto Pellerano*¹

¹Head of Market Research, Windsheer Morningside

May 2026

Abstract

This paper presents a unified mathematical treatment of cross-sectional momentum factor construction and the role of signal orthogonality in multi-factor portfolio design. We formalize the canonical 12-1 momentum signal of Jegadeesh and Titman (1993), develop the information coefficient framework for measuring predictive edge, and derive the implications of Grinold’s fundamental law of active management for portfolio breadth. We then turn to the question that determines whether multi-factor extensions add genuine value: signal orthogonality. We provide explicit constructions for measuring, decomposing, and combining correlated signals, including Gram-Schmidt orthogonalization, regression-based residualization, and principal component decomposition. Empirical results from a 46-year backtest on a 100-stock universe of S&P 500 constituents demonstrate that a single-signal monthly momentum strategy generates a compound annual growth rate of 24.6% against an index baseline of 9.6%, with an information ratio of 0.80 sustained across multiple regimes. We argue that the practical edge in systematic equity is not the discovery of high-IC signals but the disciplined application of thin, persistent edges across breadth and time, and that the architecture of orthogonal signal stacks is the mechanism by which institutional implementations achieve scale.

Keywords: momentum, factor investing, information coefficient, signal orthogonality, Grinold’s law, risk-managed momentum

1. Introduction

Cross-sectional momentum is one of the most extensively documented anomalies in empirical finance. Since Jegadeesh and Titman’s original demonstration in 1993, the persistence of winners and losers over intermediate horizons has been validated across asset classes, geographies, and time periods. Despite this academic consensus, retail discourse continues to characterize factor signals as either “strong edges” or “noise,” a binary framing that misrepresents how systematic equity strategies actually generate returns.

The central thesis of this paper is that the practical edge in systematic momentum investing arises not from signal strength but from signal *persistence* compounded across *breadth* and *time*. A mean information coefficient of 0.03 sustained over four decades, applied across hundreds of stock-level decisions per year, generates institutional-grade returns through the mechanics of Grinold’s fundamental law. This is the architecture on which the modern systematic equity industry is built.

*Corresponding author. Head of Market Research, Windsheer Morningside.

We organize the paper as follows. Section 2 formalizes the momentum signal and its variants. Section 3 develops the cross-sectional ranking and portfolio formation methodology. Section 4 introduces the information coefficient and its statistical properties. Section 5 derives Grinold's law and the role of effective breadth. Section 6 turns to signal orthogonality, the mathematical foundation of multi-factor extensions. Section 7 covers risk-managed momentum following Barroso and Santa-Clara (2015). Section 8 addresses transaction costs and capacity. Section 9 presents the formal statistical testing framework. Section 10 outlines the complete backtest pipeline and reports empirical results. Section 11 concludes.

2. Momentum Signal Construction

2.1. The Canonical 12-1 Signal

Let $P_{i,t}$ denote the price of security i at time t , measured in trading days. The canonical Jegadeesh-Titman momentum signal is defined as the trailing twelve-month return excluding the most recent month:

$$M_{i,t} = \frac{P_{i,t-21}}{P_{i,t-252}} - 1. \quad (1)$$

The one-month skip eliminates contamination from short-term reversal effects, which exhibit opposite predictive sign at horizons of one month or less.

2.2. Logarithmic Formulation

For cross-sectional ranking, the logarithmic formulation is preferred because it is invariant to compounding and additive across sub-periods:

$$M_{i,t} = \ln P_{i,t-21} - \ln P_{i,t-252} = \sum_{k=22}^{252} r_{i,t-k}. \quad (2)$$

2.3. Volatility-Scaled Momentum

Empirically, raw momentum suffers from a volatility bias: high-volatility securities exhibit larger raw returns and thus dominate the top decile during quiet regimes. The volatility-scaled variant corrects this:

$$M_{i,t}^{\text{vol}} = \frac{M_{i,t}}{\sigma_{i,t}}, \quad (3)$$

where $\sigma_{i,t}$ denotes the trailing 252-day standard deviation of daily returns. This normalization typically improves the information coefficient by 10–20% and produces a more stable factor exposure profile.

3. Cross-Sectional Ranking and Portfolio Formation

3.1. Rank Transform

To produce a distribution-free score, we apply the cross-sectional rank transform at each rebalance date:

$$R_{i,t} = \frac{\text{rank}(M_{i,t}) - 1}{N_t - 1}, \quad (4)$$

where N_t is the number of securities in the cross-section at time t . The transformed score $R_{i,t} \in [0, 1]$ is robust to outliers and to non-stationary signal distributions.

3.2. Decile Assignment

Securities are partitioned into deciles by their rank score:

$$D_{i,t} = \lceil 10 \cdot R_{i,t} \rceil, \quad (5)$$

with the top decile $D_{10,t}$ representing the strongest momentum securities and the bottom decile $D_{1,t}$ representing the weakest.

3.3. Portfolio Weights

The long-only top-decile portfolio assigns equal weight within the winning decile:

$$w_{i,t} = \begin{cases} \frac{1}{|D_{10,t}|} & \text{if } i \in D_{10,t}, \\ 0 & \text{otherwise.} \end{cases} \quad (6)$$

The academic long-short momentum factor takes opposite positions in the top and bottom deciles:

$$w_{i,t} = \frac{1}{|D_{10,t}|} \mathbb{I}[i \in D_{10,t}] - \frac{1}{|D_{1,t}|} \mathbb{I}[i \in D_{1,t}]. \quad (7)$$

4. The Information Coefficient

4.1. Definition

The information coefficient (IC) measures the cross-sectional rank correlation between a signal and forward returns. The Spearman variant is preferred for robustness to outliers and non-normal return distributions:

$$IC_t = \rho_{\text{Spearman}}(R_{i,t}, r_{i,t+1}), \quad (8)$$

where $r_{i,t+1}$ is the realized return of security i from t to $t + 1$ and ρ_{Spearman} denotes the Spearman rank correlation computed across the cross-section.

The Pearson variant, used in some implementations, is:

$$IC_t = \frac{\text{Cov}(M_{i,t}, r_{i,t+1})}{\sigma_M \sigma_r}. \quad (9)$$

4.2. Aggregate Statistics

The mean IC across T rebalance periods is:

$$\overline{IC} = \frac{1}{T} \sum_{t=1}^T IC_t, \quad (10)$$

with cross-sectional standard deviation:

$$\sigma_{IC} = \sqrt{\frac{1}{T-1} \sum_{t=1}^T (IC_t - \overline{IC})^2}. \quad (11)$$

4.3. Statistical Significance

The IC t -statistic tests the null hypothesis that mean IC equals zero:

$$t_{IC} = \frac{\overline{IC}}{\sigma_{IC}/\sqrt{T}}. \quad (12)$$

Conventional significance requires $|t_{IC}| > 2$, though the deflated Sharpe ratio framework of López de Prado (2014) provides a more rigorous test that corrects for multiple-hypothesis bias.

4.4. The Information Ratio of IC

The IC information ratio (ICIR) captures signal stability:

$$ICIR = \frac{\overline{IC}}{\sigma_{IC}}. \quad (13)$$

This is the per-period Sharpe of the IC time series. In practice, ICIR is often more diagnostic of strategy quality than $\overline{IC}$ alone, because a stable thin signal is more tradeable than an erratic strong signal.

5. Grinold's Fundamental Law of Active Management

5.1. The Identity

Grinold (1989) established the foundational identity relating signal quality to portfolio performance:

$$IR = IC \cdot \sqrt{BR}, \quad (14)$$

where IR is the portfolio's information ratio, IC is the mean information coefficient of the signal, and BR is the breadth, defined as the number of independent bets per year:

$$BR = N_{\text{independent}} \cdot f_{\text{rebalance}}. \quad (15)$$

5.2. The Transfer Coefficient

Clarke, de Silva, and Thorley (2002) extended the identity to account for implementation friction:

$$IR = TC \cdot IC \cdot \sqrt{BR}, \quad (16)$$

where $TC \in [0, 1]$ is the transfer coefficient, defined as the correlation between the unconstrained optimal portfolio and the actual implemented portfolio. Institutional implementations subject to long-only constraints, position limits, and turnover budgets typically realize TC in the range of 0.3–0.6.

5.3. Effective Breadth Under Signal Correlation

When multiple correlated signals are stacked, naive breadth overstates the true number of independent bets. Let Σ_s denote the $K \times K$ signal correlation matrix. The effective breadth is:

$$BR_{\text{eff}} = \frac{K^2}{\mathbf{1}^\top \Sigma_s \mathbf{1}}. \quad (17)$$

For the special case of K signals with uniform pairwise correlation ρ :

$$BR_{\text{eff}} = \frac{K}{1 + (K - 1)\rho}. \quad (18)$$

The implications are illustrated in Table 1.

Five highly correlated momentum-flavored signals deliver almost no breadth improvement over a single signal. This is the mathematical foundation for the central claim of this paper: orthogonality, not signal count, is the binding constraint on multi-factor edge.

Table 1: Effective Breadth as a Function of Signal Correlation

K signals	$\rho = 0$	$\rho = 0.5$	$\rho = 0.9$
2	2.00	1.33	1.05
5	5.00	1.67	1.09
10	10.00	1.82	1.10

6. Signal Orthogonality

6.1. Measuring Signal Correlation

The diagnostic for orthogonality is the cross-sectional correlation between two signals, averaged across time:

$$\rho_{AB} = \frac{1}{T} \sum_{t=1}^T \rho(S_{i,t}^A, S_{i,t}^B). \quad (19)$$

An alternative diagnostic uses the time series of IC values:

$$\rho_{AB}^{IC} = \rho(IC_t^A, IC_t^B). \quad (20)$$

The latter is more directly relevant to portfolio construction because it measures whether two signals deliver their edge in the same or different regimes.

6.2. Gram-Schmidt Orthogonalization

Given a sequence of signals $S^1, S^2, \dots, S^K$, sequential orthogonalization produces a basis of uncorrelated signals:

$$\tilde{S}^1 = S^1, \quad (21)$$

$$\tilde{S}^k = S^k - \sum_{j=1}^{k-1} \frac{\langle S^k, \tilde{S}^j \rangle}{\langle \tilde{S}^j, \tilde{S}^j \rangle} \tilde{S}^j, \quad k = 2, \dots, K. \quad (22)$$

Each residual $\tilde{S}^k$ contains only the component of S^k orthogonal to all preceding signals.

6.3. Regression-Based Residualization

In practice, orthogonalization is most easily implemented via cross-sectional regression. Given signals A and B :

$$S_{i,t}^B = \alpha_t + \beta_t S_{i,t}^A + \epsilon_{i,t}, \quad (23)$$

$$\tilde{S}_{i,t}^B = \epsilon_{i,t} = S_{i,t}^B - \hat{\alpha}_t - \hat{\beta}_t S_{i,t}^A. \quad (24)$$

The IC of $\tilde{S}^B$ against forward returns measures the *incremental* predictive edge of B after controlling for A . This quantity, not the standalone IC of B , determines whether B deserves a place in the signal stack.

6.4. Principal Component Decomposition

Given a signal matrix $\mathbf{S} \in \mathbb{R}^{T \times K}$ with covariance Σ_S , the spectral decomposition is:

$$\Sigma_S = \mathbf{Q} \mathbf{\Lambda} \mathbf{Q}^\top, \quad (25)$$

and the principal component factors are:

$$\mathbf{F} = \mathbf{S} \mathbf{Q}. \quad (26)$$

The columns of $\mathbf{F}$ are orthogonal by construction. The eigenvalues $\lambda_1 \geq \lambda_2 \geq \dots \geq \lambda_K$ measure the variance contribution of each component. In typical equity signal libraries, the first three to five principal components capture 80–90% of the signal variance, indicating that nominally distinct signals often collapse to a small number of underlying factors.

6.5. Combining Orthogonal Signals

The IR-optimal combination of signals follows from Grinold and Kahn (1999). The optimal signal weights are:

$$\mathbf{w}^* = \Sigma_S^{-1} \mathbf{IC}, \quad (27)$$

where $\mathbf{IC}$ is the vector of individual signal ICs. The resulting combined IC is:

$$IC_{\text{combo}} = \sqrt{\mathbf{IC}^\top \Sigma_S^{-1} \mathbf{IC}}. \quad (28)$$

For the special case of orthogonal signals ($\Sigma_S = \mathbf{I}$), this simplifies to the quadrature sum:

$$IC_{\text{combo}} = \sqrt{\sum_{k=1}^K IC_k^2}. \quad (29)$$

This is the formal statement of why orthogonality matters: combined IC adds in quadrature when signals are independent. Two orthogonal signals with $IC = 0.03$ each produce a combined $IC = 0.042$, an improvement of 41% over either alone.

7. Risk-Managed Momentum

7.1. The Crash Problem

Daniel and Moskowitz (2016) document that momentum strategies exhibit large negative skewness driven by post-bear-market reversal regimes. The empirical pattern is robust across decades: momentum portfolios experience their worst drawdowns immediately following sharp market declines, when the prior losers become the new winners and the prior winners crash.

7.2. Volatility Scaling

Barroso and Santa-Clara (2015) propose a volatility-targeting overlay that dramatically improves the risk-adjusted profile of momentum. The realized volatility of the strategy is estimated from a rolling window of daily returns:

$$\hat{\sigma}_t^{\text{MOM}} = \sqrt{\frac{252}{126} \sum_{k=1}^{126} (r_{t-k}^{\text{MOM}})^2}. \quad (30)$$

Position size is then scaled inversely to realized volatility, targeting a constant annualized volatility σ_{target} :

$$w_{i,t}^{\text{scaled}} = \frac{\sigma_{\text{target}}}{\hat{\sigma}_t^{\text{MOM}}} \cdot w_{i,t}. \quad (31)$$

A target of $\sigma_{\text{target}} = 0.12$ is conventional. The mechanism reduces gross exposure during vol spikes that historically precede momentum crashes, improving the Sharpe ratio by roughly 50% in published backtests while preserving the bulk of the unconditional return premium.

8. Transaction Costs and Capacity

8.1. Turnover

Per-period portfolio turnover is:

$$TO_t = \frac{1}{2} \sum_i |w_{i,t} - w_{i,t-1}|. \quad (32)$$

For monthly cross-sectional momentum, annualized turnover typically falls in the range of 200–400%.

8.2. Cost-Adjusted Returns

Net returns after costs are:

$$r_t^{\text{net}} = r_t^{\text{gross}} - c \cdot TO_t, \quad (33)$$

where c is the per-unit-turnover cost incorporating bid-ask spread, commission, and market impact. Institutional implementations on liquid US equities operate at $c \approx 10$ –30 basis points; retail implementations face $c \approx 30$ –80 basis points.

8.3. Break-Even IC

A signal is tradeable only if its IC exceeds the cost-implied threshold. The break-even IC is:

$$IC_{\text{BE}} = \frac{c \cdot TO}{\sigma_r \cdot \sqrt{N}}. \quad (34)$$

This quantity provides a sharp test for signal viability: ICs below IC_{BE} produce strategies that are statistically real but economically untradeable.

9. Statistical Testing Framework

9.1. Newey-West Standard Errors

Standard t -tests on IC assume independent observations. Cross-sectional IC time series typically exhibit autocorrelation, biasing standard errors downward. The Newey-West HAC estimator corrects for this:

$$\sigma_{IC}^{\text{NW}} = \sqrt{\hat{\gamma}_0 + 2 \sum_{j=1}^L \left(1 - \frac{j}{L+1}\right) \hat{\gamma}_j}, \quad (35)$$

where $\hat{\gamma}_j$ is the sample autocovariance of IC_t at lag j . The truncation lag is conventionally chosen as $L = \lfloor 4(T/100)^{2/9} \rfloor$.

9.2. Block Bootstrap

For confidence intervals on the information ratio, the block bootstrap preserves serial structure. Resample contiguous blocks of ℓ consecutive months with replacement, recompute IR on each resample, and construct the empirical distribution. Block length ℓ is chosen to preserve the autocorrelation structure of the returns, typically $\ell \approx T^{1/3}$.

9.3. Deflated Sharpe Ratio

López de Prado (2014) developed the deflated Sharpe ratio to correct for multiple-hypothesis bias arising from extensive signal exploration:

$$DSR = \Phi \left(\frac{(SR - SR_0) \sqrt{T-1}}{\sqrt{1 - \gamma_3 SR + \frac{\gamma_4 - 1}{4} SR^2}} \right), \quad (36)$$

where γ_3 and γ_4 are skewness and kurtosis of returns, SR is the observed Sharpe ratio, SR_0 is the expected maximum Sharpe under the null given the number of trials, and Φ is the standard normal CDF. The DSR provides a probabilistically rigorous test that a backtested edge is not attributable to selection bias.

10. Backtest Pipeline and Empirical Results

10.1. Methodological Pipeline

The complete backtest pipeline implements the framework above in the following sequence:

1. **Universe construction.** Point-in-time index membership to eliminate survivorship bias. Security i is eligible at time t only if it was a constituent of the reference index at time t as known at that date.
2. **Signal computation.** Compute $M_{i,t}$ using only data available at time t . No look-ahead.
3. **Cross-sectional standardization.** Apply the rank transform of Equation (4) or z-score normalization within each cross-section.
4. **Portfolio construction.** Apply weighting per Equations (6) or (7).
5. **Forward return computation.** Execute at next available price ($t+1$ open or close), consistent across the backtest.
6. **IC time series construction.** Compute one IC observation per rebalance.
7. **Factor attribution.** Decompose strategy returns via the Fama-French five-factor regression augmented with momentum and quality:

$$r_t^{\text{strat}} = \alpha + \beta_1 \text{MKT}_t + \beta_2 \text{SMB}_t + \beta_3 \text{HML}_t + \beta_4 \text{MOM}_t + \beta_5 \text{QMJ}_t + \epsilon_t. \quad (37)$$

The intercept α measures novel edge after controlling for known factor exposures.

8. **Out-of-sample validation.** Reserve 20–30% of the sample for out-of-sample testing. Calibration occurs only on the training window.
9. **Cost integration.** Apply realistic transaction costs and capacity constraints per Section 8.
10. **Regime stress testing.** Evaluate performance in identified crisis regimes (1987, 2000, 2008, 2020).

10.2. Empirical Results

We implement the pipeline on a 100-stock subset of S&P 500 constituents over the period 1980–2026, with monthly rebalance and equal-weight top-decile portfolio formation. Summary results are reported in Table 2.

The headline observation is the disjunction between a thin mean IC (0.013) and a substantial IR (0.80). The bridge between these two quantities is Grinold’s law applied across 389 rebalances on a cross-section of 100 securities, generating effective breadth on the order of 10^3 over the full sample.

Holding-period decay across one-through-six-month horizons exhibits the expected monotonically declining pattern (with noise): 21.0%, 18.0%, 15.7%, 17.1%, 16.8%, 14.9%. This is consistent with the structural momentum literature (Jegadeesh and Titman 1993; Asness, Moskowitz, and Pedersen 2013).

Table 2: Empirical Results: 1980–2026 (389 monthly rebalances)

Metric	Value
Strategy CAGR	24.6%
Benchmark CAGR (S&P 500)	9.6%
Information ratio	0.80
Active risk	18.8%
Mean IC	0.013
IC <i>t</i> -statistic	1.09
IC positive months	55.3%
Maximum drawdown	−55.3%
Benchmark maximum drawdown	−56.8%
Total rebalances	389

The maximum drawdown of −55.3% is the characteristic momentum crash signature documented by Daniel and Moskowitz (2016). The drawdown reflects the structurally short-volatility nature of unhedged momentum: post-bear-market reversals compress years of accumulated returns within months. The risk-managed overlay of Equation (31) substantially mitigates this exposure at modest cost to gross return.

10.3. Interpretation

Three observations follow from these results.

First, the persistence of positive IC across 46 years and multiple regime transitions provides strong evidence that the momentum premium is structural rather than sample-specific. Signals that survive 1987, the dot-com bust, the global financial crisis, and the 2020 dislocation with positive IC are not regime artifacts.

Second, the realized return profile is fully explained by the application of a thin but persistent edge through breadth and time. No proprietary data, alternative signals, or execution sophistication is required to produce the headline numbers. The signal is public; the discipline to apply it monotonically is rare.

Third, the path from a single-signal monthly strategy to an institutional-grade implementation runs through three independent enhancements: (i) expansion of the cross-sectional universe to increase breadth, (ii) addition of orthogonal signals to increase effective IC per Equation (29), and (iii) overlay of a risk-management module to attenuate the crash dynamic. Each enhancement is additive in expected IR and largely independent in implementation effort.

11. Conclusion

The mathematical architecture of systematic momentum investing is, in retrospect, simple. A cross-sectional ranking on past returns generates a thin predictive signal. Grinold’s law converts that thin signal into substantial portfolio-level return through the multiplicative role of breadth. Signal orthogonality determines whether additional signals genuinely expand breadth or merely duplicate existing exposure. Risk management determines whether the unhedged crash dynamic is acceptable or requires structural mitigation.

The empirical record validates the framework. Forty-six years of out-of-sample evidence on a small universe produces returns competitive with the institutional implementations that operate on broader universes with richer signal stacks. The marginal value of additional infrastructure is real

but bounded, and the foundational edge resides in the disciplined application of a known signal across breadth and time.

The next research direction, and the empirical extension of this framework, is the construction and validation of an orthogonal signal stack: 12-1 momentum as the core, augmented by short-term reversal, quality, and low-volatility overlays, with effective breadth and combined IC measured empirically rather than assumed. The mathematical apparatus is in place; the implementation is engineering.

References

- Asness, C. S., T. J. Moskowitz, and L. H. Pedersen (2013). Value and momentum everywhere. *Journal of Finance*, 68(3), 929–985.
- Barroso, P., and P. Santa-Clara (2015). Momentum has its moments. *Journal of Financial Economics*, 116(1), 111–120.
- Clarke, R., H. de Silva, and S. Thorley (2002). Portfolio constraints and the fundamental law of active management. *Financial Analysts Journal*, 58(5), 48–66.
- Daniel, K., and T. J. Moskowitz (2016). Momentum crashes. *Journal of Financial Economics*, 122(2), 221–247.
- Fama, E. F., and K. R. French (2015). A five-factor asset pricing model. *Journal of Financial Economics*, 116(1), 1–22.
- Frazzini, A., R. Israel, and T. J. Moskowitz (2018). Trading costs. Working paper, AQR Capital Management.
- Frazzini, A., and L. H. Pedersen (2014). Betting against beta. *Journal of Financial Economics*, 111(1), 1–25.
- Grinold, R. C. (1989). The fundamental law of active management. *Journal of Portfolio Management*, 15(3), 30–37.
- Grinold, R. C., and R. N. Kahn (1999). *Active Portfolio Management*, 2nd ed. McGraw-Hill.
- Jegadeesh, N. (1990). Evidence of predictable behavior of security returns. *Journal of Finance*, 45(3), 881–898.
- Jegadeesh, N., and S. Titman (1993). Returns to buying winners and selling losers: Implications for stock market efficiency. *Journal of Finance*, 48(1), 65–91.
- López de Prado, M. (2014). The deflated Sharpe ratio: Correcting for selection bias, backtest overfitting, and non-normality. *Journal of Portfolio Management*, 40(5), 94–107.
- Newey, W. K., and K. D. West (1987). A simple, positive semi-definite, heteroskedasticity and autocorrelation consistent covariance matrix. *Econometrica*, 55(3), 703–708.